

Teach a Quadcopter How to Fly

REVIEW

CODE REVIEW 1

HISTORY

Meets Specifications

Hello,

You've done a great work by implementing the DDPG and to teach quadcopter how to fly. Your network architecture is quite well designed and similar to the original DDPG paper, which also helped in better performance. Awesome job by providing the nice visualizations.

Also, noise is only for the purpose of exploration. At the time of testing the agent, you should not add the noise.

Since you have completed this project, I would recommend you to know how to apply these techniques to real-world problems. For that, I would advise you to go through [Deep Reinforcement Learning for Self Driving Car by MIT](#). So that you can know more about reinforcement learning algorithms in a broader and real-world perspective.

Define the Task, Define the Agent, and Train Your Agent!

The `agent.py` file contains a functional implementation of a reinforcement learning algorithm.

Awesome

- Successful implementation of the DDPG algorithm.
- Actor and critic network are correctly implemented.
- It was good to use a target network for both actor and critic network as suggested by the author of DDPG Paper.
- Further using soft updates for the target network was a good choice.

The `Quadcopter_Project.ipynb` notebook includes code to train the agent.

Plot the Rewards

A plot of rewards per episode is used to illustrate how the agent learns over time.

Reflections

The submission describes the task and reward function, and the description lines up with the implementation in `task.py`. It is clear how the reward function can be used to guide the agent to accomplish the task.

Awesome

- Task and Reward function are explained properly.

The submission provides a detailed description of the agent in `agent.py`.

Awesome

- Good job by choosing the DDPG algorithm for continuous state space problem.
- Good job by including hyperparameters and network architecture in the notebook.

The submission discusses the rewards plot. Ideally, the plot shows that the agent has learned (with episode rewards that are gradually increasing). If not, the submission describes in detail various attempted settings (hyperparameters and architectures, etc) that were tested to teach the agent.

Awesome

- It is quite common to see agent not learning ideally in `Quadcopter` project since it is not an easy project. And you have done a good job so far.

A brief overall summary of the experience working on the project is provided, with ideas for further improving the project.



RETURN TO PATH
